

SME Operating Agent Business Proposal

A practical agent for company knowledge meeting follow up and founder led recruiting

Business Proposal Draft

Team code	Pending confirmation
Submission	Show Me Your Agents Hackathon
Repository	github.com/joshualumzy/personal_agent_long_context
Date	27 September 2026

Proposal purpose

Secure approval to pilot one integrated agent that helps a small team recover company context, convert meeting commitments into controlled follow up, and run a structured recruiting workflow.

Executive Summary

Small companies often rely on the same people to answer internal questions, remember meeting commitments, and run hiring. The information needed for those tasks is spread across messages, documents, tickets, meeting transcripts, and personal memory. Staff must reconstruct the context, prepare the work, and remember to follow through while continuing their main jobs.

We propose the SME Operating Agent, a single product with three connected workflows. Company Context retrieves employee-visible evidence and returns inspectable citations. Meeting Actions converts text transcript lines into traceable answers, conflicts, and reviewable drafts. Recruiting helps a founder define criteria, search and assess public professional profiles, and prepare outreach without sending on the founder's behalf.

The current prototype demonstrates the product architecture and human-control model. It does not yet establish business savings or production reliability. The recommended next step is a bounded SME pilot that measures time saved, follow-through, quality, and control failures against the team's current workflow.

One Business Story

Know	Act	Grow
Recover company context from inspectable evidence	Turn text meetings into controlled follow up work	Run a structured founder led recruiting workflow

These workflows share one operating principle: the agent may retrieve, analyse, and draft, but important external actions remain visible to a person. Company questions use cited evidence. Meeting drafts require the correct policy tier and exact-payload approval. Recruiting outreach stays in the founder's own Gmail or LinkedIn account.

Recommendation

Proceed to a four to six week pilot with one SME team, one recurring meeting type, and one active hiring role. Establish baselines before activation, configure only the integrations needed for the pilot, and treat the proposed KPI targets in this document as acceptance hypotheses rather than achieved results.

Problem and Opportunity

Target Users

The primary users are employees and founders in small and medium enterprises who must turn scattered company information into daily work without dedicated knowledge-management, project-operations, or recruiting capacity. Managers, IT administrators, and governance owners are secondary stakeholders because they need reliable sources, clear permissions, and visible accountability.

Current Workflow

1. An employee receives a company question, makes a commitment in a meeting, or a founder decides to hire.
2. They search several systems and reconstruct the relevant history by hand.
3. They translate what they found into an answer, email, ticket, calendar event, document, or hiring criteria.
4. They chase missing details and remember which action still needs follow up.
5. They personally check and execute the consequential action.

Why the Problem Matters

Operational problem	Business consequence	Evidence status
Fragmented company records	Slower answers and decisions based on incomplete context	Demonstrated on a synthetic company corpus; real SME baseline required
Meeting commitments rely on memory	Follow up is delayed, duplicated, or missed	Implemented workflow; real-team frequency and loss rate required
Founder led hiring has little support	Search, screening, and follow up compete with core leadership work	Implemented workflow; real hiring-time baseline required

The opportunity is to reduce the manual gap between finding context and preparing the next piece of work. The proposal does not assume that every task should be automated. It focuses on the repetitive reconstruction, drafting, and tracking steps while retaining human judgment for decisions and external effects.

Problem Validation Needed

- Interview employees and founders about the frequency and consequence of these tasks.
- Measure current time, manual steps, missed follow ups, and rework for representative cases.
- Document the systems, permissions, and approval practices already used by the pilot SME.

Proposed Solution

Company Context

An authenticated employee asks a company question. The agent searches employee-visible Company Evidence using full-text retrieval, vector similarity, and explicit links between related records. It returns a concise answer with inspectable source identifiers or states that the available evidence is insufficient. Employee-specific Personal Memory is displayed separately and is not treated as company evidence.

Meeting Actions

The meeting workflow accepts transcript text entered live, pasted, or replayed from a scripted or OrgForge meeting. It does not perform speech-to-text transcription. Each segment is screened before model use. The service identifies commitments, questions, and decisions, keeps the exact trigger quote, answers company questions through Company Context, checks for conflicting decisions, and prepares supported drafts.

Policy tier	Examples	Control
Automatic	Read-only answer or conflict flag	Executed and traced without changing an external system
Approval	Email, message, calendar, ticket, document, spreadsheet, hiring request	Approval is bound to the exact payload; an edit requires approval again
Escalation	Money or contract commitment	The employee cannot execute it; the action is raised to a named authority
Blocked	Secret or transcript instruction attack	Stopped before model use

Recruiting

A founder states a hiring need or uploads a job description. The agent proposes reviewable criteria, searches public professional profiles through Exa when configured or uses clearly labelled fictional sample profiles, and scores candidates against the confirmed criteria. Founder feedback can change criteria or create a proposal, but the change applies only after an explicit decision. The agent may find a work email through configured lookup services and draft outreach; it never guesses an address and cannot send through its chat tools.

How the Workflows Connect

- Company Context supplies evidence and answers to Meeting Actions.
- An approved hiring need from a meeting creates a draft role in Recruiting.
- Recruiting is also available from the main company chat as a skill that loads only for hiring requests.
- Personal Memory retains employee context and stated hiring rationale through separate adapters, while Company Evidence remains in PostgreSQL.

Business Value

The business case depends on measurable reductions in manual work and avoidable rework. The prototype supports the workflow, but the following benefits remain hypotheses until a pilot measures them against a real SME baseline.

Value area	How value is created	Pilot evidence required
Productivity	Fewer manual searches and less time converting context into drafts	Median handling time and manual steps before and after
Follow through	Meeting commitments become visible, traceable work instead of relying on memory	Completion rate, time to first draft, and overdue actions
Hiring capacity	A founder receives structured criteria, a reviewable pool, and follow up drafts	Founder hours per role and time from requirement to reviewed pool
Service consistency	Answers use inspectable evidence and unsupported questions can stop explicitly	Correctness, evidence recall, and user source inspection
Risk control	Deterministic policy, exact-payload approval, and user-account handoffs limit silent action	Approval bypasses, stale approvals, unsafe content, and outbound-send violations
Scale	The same evidence and human-control foundation supports several workflows	Users, cases, latency, operating cost, and support effort

Before and After

Task	Current process	Proposed process
Company question	Search systems, connect records, write an answer	Retrieve evidence, inspect citations, accept the answer or insufficient-evidence result
Meeting commitment	Take notes, interpret the promise, gather details, prepare work later	Screen transcript text, identify the commitment, draft now, review the exact payload, hand off
Founder hiring	Write criteria, search profiles, compare candidates, remember feedback, chase replies	Confirm criteria, review a scored pool, record feedback, prepare outreach and follow ups

Commercial Position

No revenue, cost-saving, or return-on-investment claim is established today. A pilot should first determine whether the saved staff and founder time exceeds the operating cost of model calls, storage, integrations, onboarding, and support. Pricing and packaging should follow that evidence rather than precede it.

Impact and Outcomes

The pilot should compare representative cases before and after the agent is introduced. Proposed targets below are decision thresholds for the pilot, not results already achieved. Baselines must be measured during the first week and targets adjusted if the task mix differs materially from the assumptions.

Outcome	Metric	Proposed pilot target
Faster context reconstruction	Median time to produce an evidence-backed answer	At least 30 percent lower than baseline
Less meeting administration	Median time from meeting end to reviewable follow up draft	At least 50 percent lower than baseline
Better follow through	Share of accepted meeting actions completed by their due date	At least 90 percent
Lower founder hiring effort	Founder administration time per active role	At least 25 percent lower than baseline
Trustworthy company answers	Factual accuracy and expected-evidence recall on pilot questions	Threshold set after baseline; both must improve from the current prototype
Human control	Unauthorized outbound sends or executed escalated actions	Zero
Usable evidence	Pilot users who can open and understand the source behind an answer	At least 90 percent

Measurement Plan

1. Capture baseline duration, steps, rework, and completion for a representative set of company questions, meetings, and one hiring role.
2. Run the same categories through the agent and log retrieval, model, policy, approval, handoff, error, and completion timestamps.
3. Review answer correctness and source support with an authorised SME owner; review hiring judgments for relevance and fairness with a human hiring owner.
4. Conduct weekly user interviews and inspect failure cases, not only aggregate scores.
5. At the end of the pilot, compare value, risk, operating cost, and support effort before deciding whether to expand.

Current Technical Evidence

Area	Recorded result	Interpretation
Company Context	76-question run: 12 percent factual accuracy, 29 percent expected-evidence recall, 100 percent citation-ID integrity, 39.68 seconds mean latency	Citation identifiers were controlled in that sample, but answer quality and latency require improvement
Meeting Actions	One live run: 23 of 23 policy tiers correct; 9 of 9 unsafe lines blocked; 0 of 44 ordinary lines blocked; 1 of 1 conflict found; email recall 2 of 3	Encouraging bounded evidence from one model run, not a reliability guarantee
Recruiting	Deterministic, holdout, regression, chaos, and browser test assets exist	A final submission-snapshot quality summary remains to be produced

These results support a pilot recommendation, not a production claim. The company-answer evaluation is the clearest current quality gap. The final submission should show improvement or explain why the pilot is designed to address it.

Feasibility and Scalability

What Exists Today

- A TypeScript and Fastify modular monolith with three registered browser surfaces.
- PostgreSQL and pgvector for Company Evidence, conversations, employee records, meeting state, and action logs.
- A multi-step company reasoning loop with source validation and Server-Sent Events.
- A meeting service with pre-model screening, extraction, evidence gathering, deterministic policy, payload versioning, approval, and handoffs.
- A per-role recruiting service with criteria, search, scoring, feedback, proposals, outreach drafts, reply handling, retention, and a main-chat skill.
- Letta adapters for employee context and founder hiring-intent memory, separate from Company Evidence.

Operating Requirements

Capability	Required for core pilot	Optional live integration
Application	Node.js runtime, Fastify server, session secret	Public deployment and managed secrets
Company evidence	PostgreSQL with pgvector and permission-scoped data	Incremental connectors for approved systems
Models and memory	SoCLaaS or supported model path; Letta or documented fallback	Provider redundancy and monitored service levels
Recruiting	Labelled sample profiles or Exa; manual outreach completion	Exa, Hunter, Prospeo, Google OAuth, read-only LinkedIn sync
Meeting handoffs	Browser links and explicit simulated state where required	Configured GitHub repository and approved office-suite choices

Path to Scale

1. Prototype to pilot: select one SME, one department, one meeting type, and one active hiring role.
2. Pilot hardening: add permission rules, consent and retention controls, monitoring, support procedures, and a bounded evaluation set.
3. Production transition: replace demo personas and local recruiting files with tenant-scoped identity, authorisation, and durable storage.
4. Controlled expansion: add connector-specific ingestion and approval contracts only after each workflow meets its quality and safety threshold.
5. Cross-functional scale: reuse the shared evidence and human-control foundation for other departments while preserving separate permissions and measures.

Scalability Questions to Resolve

- Expected user count, case volume, latency requirement, and model budget
- Permission model across departments and data sources

- Operational ownership for monitoring, incident response, backup, and deletion
- Legal basis and fairness controls for recruiting data and decisions

Risk Governance and Human Control

Risk	Current control	Pilot requirement
Incorrect answer with valid-looking sources	Inspectable evidence and citation-ID validation	Human-labelled answer and semantic source-support review
Missing evidence	Full-text, vector, and explicit-link retrieval; insufficient-evidence outcome	Improve recall against pilot questions before expansion
Unsafe meeting content	Pre-model screening for prohibited data and instruction attacks	Adversarial testing with SME-specific examples
Stale or unauthorised action	Deterministic tier and exact-payload approval; edits invalidate approval	Server-derived identity, audit review, and negative approval tests
Silent external send	S2 and S3 prepare handoffs; S3 chat tools expose no send operation	Confirm zero outbound sends without a user's action
Recruiting privacy or bias	Public professional fields, no guessed email, private pass reasons blocked, closed-candidate retention rule	Legal review, fairness measures, consent basis, and deletion verification
Evaluation leakage	Reserved oracle records excluded from the runtime Company Evidence path	Re-verify the final deployed dataset and access boundary

Claim Boundary

The proposal may claim that the current code integrates three workflows and implements the stated controls. It may report saved evaluation results with their sample and run limitations. It should not claim production readiness, validation on real SME data, complete prompt-injection resistance, unbiased hiring, autonomous transcription, autonomous sending, or proven financial returns.

Pilot Decision

Approve a bounded pilot if the SME can provide an authorised evidence set, a recurring meeting workflow, one active hiring role, and owners for security and outcome review. Do not expand beyond the pilot until answer quality improves, human-control tests pass, recruiting governance is agreed, and operating cost is measured.

Information Required Before Submission

Item	Why it matters
Team code and final product name	Required submission identity
Real SME problem evidence	Supports the opportunity and current-workflow claims
Final deployment URL and evidence	Shows that the prototype can be accessed and operated
Final verification and S3 evaluation summary	Aligns technical evidence with the proposal and demo
Demo scenario and configured integrations	Ensures the proposal and recorded demonstration tell the same story

Conclusion

The SME Operating Agent addresses a connected operating problem: small teams have to recover context, convert decisions into work, and sustain hiring with limited specialist support. The current implementation is credible enough for a controlled pilot because it joins the three workflows and keeps consequential actions visible to people. The business case now depends on validating the problem with an SME, improving company-answer quality, and measuring whether the saved effort justifies the operating cost and governance burden.