

SME Operating Agent

Business Proposal

Know the company. Act on meetings. Grow the team.

Company	Stellar Ark AI
Team code	Pending confirmation
Hackathon	Show Me Your Agents
Repository	github.com/joshualumzy/personal_agent_long_context
Date	27 September 2026

Proposal objective

Approve a bounded internal pilot of one integrated agent that helps Stellar Ark AI employees recover company context, convert meeting decisions into controlled follow up, and support founder-led recruiting.

Executive Summary

Stellar Ark AI builds AI-native wearables, including context glasses and a pendant, paired with a proactive agent system that turns everyday conversations into persistent memory and action. Its agent layer connects to Feishu, email, calendar, and CRM for customers in consulting, finance, legal, enterprise sales, and technology.

The company faces a related internal challenge: employees must recover knowledge from scattered records, translate meeting decisions into accountable work, and support hiring while continuing core product and customer responsibilities. These are connected failures in the path from context to action.

The SME Operating Agent joins three working surfaces into one business story. Company Context provides evidence-backed answers. Meeting Actions turns text transcripts into traceable decisions and controlled follow up. Recruiting helps a founder define criteria, review candidates, and prepare outreach. Important external actions remain visible to a person.

One Business Story

Know	Act	Grow
Recover company context from inspectable evidence	Turn meeting decisions into controlled follow up	Run structured founder-led recruiting

Business Ask

Run a four-to-six-week internal pilot with one Stellar Ark AI team, one recurring meeting type, and one active hiring role. Establish the baseline first, introduce only the required integrations, and decide whether to expand based on measured value, quality, risk, and operating cost.

Problem and Opportunity

Who Experiences the Problem

The primary users are Stellar Ark AI employees and founders who must turn fragmented company information into daily work inside a fast-moving AI-native product business. Managers and governance owners also need reliable sources, clear permissions, and visible accountability.

Current Workflow and Consequence

Current task	Manual friction	Business consequence
Answer a company question	Search several systems and reconstruct the history	Slower answers and decisions based on incomplete context
Follow up a meeting decision	Interpret notes, gather details, draft work, and remember ownership	Actions may be delayed, duplicated, or missed
Run founder-led hiring	Define criteria, search profiles, compare candidates, and chase replies	Recruiting competes with product, customer, and leadership work

Why This Opportunity Matters

- The same scarce employees repeatedly reconstruct context before they can act.
- Meeting follow through and hiring are high-value workflows where visibility and human judgment matter.
- A shared evidence and control layer can support several workflows without silently automating consequential decisions.

Problem Basis and Data Approach

The problem is grounded in the team's first-hand Stellar Ark AI context. Private operational records, customer information, and conversations are intentionally excluded from the hackathon. OrgForge supplies privacy-safe synthetic demo and evaluation data; it is not a customer, dependency, deployment target, or integration. The internal pilot will establish the quantitative baseline.

Proposed Solution

Workflow	What the agent does	Human control
Company Context	Retrieves employee-visible evidence using full text, vector similarity, and explicit links; answers with inspectable source identifiers or states insufficient evidence	Personal Memory stays separate and is never cited as Company Evidence
Meeting Actions	Screens text transcript lines; extracts commitments, questions, and decisions; checks evidence and conflicts; prepares supported drafts	Read-only results may complete automatically; drafts require exact-payload approval; sensitive or unauthorised actions are blocked or escalated
Recruiting	Turns a hiring need into criteria, searches and scores public profiles, records explicit feedback, and prepares outreach and follow ups	The founder confirms changes and completes outreach in their own account; the chat tools cannot send

Connected Demonstration Story

1. An employee asks a company question and receives an evidence-backed answer.
2. A meeting transcript introduces a commitment, question, decision, and hiring need.
3. Meeting Actions uses Company Context, detects conflicts, and prepares controlled follow up.
4. After approval, the hiring need becomes a draft role in Recruiting.
5. The founder reviews criteria and candidates, then completes outreach through a user-controlled handoff.

Scope Boundary

The prototype consumes text transcripts; it does not perform speech-to-text. It prepares handoffs rather than silently sending messages or changing third-party systems. These boundaries keep the hackathon story aligned with the current implementation.

Business Value and Outcomes

Value area	Expected value	How the pilot will test it
Productivity	Less time reconstructing context and converting it into work	Compare median handling time and manual steps before and after
Follow through	Visible, traceable commitments instead of reliance on memory	Measure time to draft, completion rate, and overdue actions
Hiring capacity	A structured candidate and follow-up workflow for the founder	Measure founder administration time per active role
Service quality	Inspectable evidence and an explicit insufficient-evidence outcome	Review correctness, expected-evidence recall, and source usability
Risk control	Deterministic approval, escalation, blocking, and user-controlled handoffs	Count approval bypasses, unsafe content failures, and unauthorised sends

Proposed Pilot Targets

Outcome	Decision threshold
Faster context reconstruction	At least 30 percent below baseline median time
Less meeting administration	At least 50 percent below baseline time to a reviewable draft
Better follow through	At least 90 percent of accepted actions completed by their due date
Lower founder hiring effort	At least 25 percent below baseline administration time
Human control	Zero unauthorised outbound sends or executed escalated actions

These are proposed decision thresholds, not achieved results. Week one establishes the baseline; the pilot then compares representative cases and reviews failures as well as aggregate scores.

Feasibility and Scalability

What Exists Today

- A TypeScript and Fastify modular monolith with Company Context, Meeting Actions, and Recruiting browser surfaces.
- PostgreSQL and pgvector for company evidence, conversations, meeting state, and action logs.
- Source validation, Server-Sent Events, pre-model screening, deterministic policy, exact-payload approval, and user-controlled handoffs.
- A recruiting service with criteria, search, scoring, feedback, proposals, outreach drafts, retention, and a main-chat skill.
- Letta adapters for employee and founder context, kept separate from Company Evidence.

Current Technical Evidence

Area	Recorded result	Business interpretation
Company Context	76-question run: 12% factual accuracy, 29% expected-evidence recall, 100% citation-ID integrity, 39.68s mean latency	Citation handling is controlled, but answer quality and latency are the clearest pilot-readiness gap
Meeting Actions	One live run: 23/23 policy tiers, 9/9 unsafe lines blocked, 0/44 ordinary lines falsely blocked, 1/1 conflict found	Encouraging bounded evidence from one model run, not a reliability guarantee
Recruiting	Deterministic, holdout, regression, chaos, and browser test assets exist	A concise final submission-snapshot quality result remains required

Path to Scale

1. Pilot: one team, one meeting type, one active role, and an authorised privacy-safe evidence set.
2. Harden: add tenant-scoped identity, permissions, retention, monitoring, support procedures, and bounded evaluations.
3. Expand: add connector-specific ingestion and approval contracts only after each workflow meets its quality and safety threshold.

Pilot Proposal

Pilot Design

Element	Proposal
Duration	Four to six weeks
Participants	One Stellar Ark AI team with named business, security, and outcome owners
Scope	One recurring meeting workflow, representative company questions, and one active hiring role
Data	Approved privacy-safe evidence first; private company data only after permissions, retention, and approval practices are documented
Decision	Expand only if value thresholds are met, answer quality improves, human-control tests pass, and operating cost is acceptable

Critical Controls

Risk	Pilot control
Incorrect or unsupported answers	Human-labelled review of correctness and semantic source support; preserve insufficient-evidence outcomes
Unauthorised action	Server-derived identity, deterministic policy, exact-payload approval, audit review, and negative tests
Recruiting privacy or bias	Public professional fields, no guessed email, human review, fairness measures, consent basis, and deletion verification

Decision Requested

Approve the bounded internal pilot and its measurement plan. The current prototype is credible enough to test because it connects the three workflows and keeps consequential actions visible to people. It is not yet a production system: the pilot must quantify the baseline, improve company-answer quality, verify controls, and determine whether saved effort justifies operating cost and governance burden.